

# Leveraging Xception Model for Image Classification Tasks

## Abstract

Deep learning has revolutionized image classification with advancements in convolutional neural networks (CNNs). This paper demonstrates the development and training of an Xception-based model for multi-class image classification. The model leverages transfer learning from a pre-trained Xception base and incorporates fine-tuning to enhance performance. Experimental results reveal the effectiveness of this approach in achieving high accuracy.

---

## 1. Introduction

Deep learning architectures have proven their capability in solving complex image classification problems. Among them, the Xception model stands out due to its use of depthwise separable convolutions, which reduce computational overhead while maintaining high accuracy. This paper details the construction and training of an Xception model tailored for a custom image classification task.

## Methodology

### 2.1 Xception Architecture

The Xception model is based on an "extreme" version of the Inception architecture, where traditional inception modules are replaced with depthwise separable convolutions. This makes it computationally efficient while maintaining state-of-the-art performance.

### 2.2 Model Construction

The following Python function outlines the steps to build the Xception model:

```
def build_xception_model():
    base_model = Xception(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(512, activation='relu')(x)
    predictions = Dense(len(classes), activation='softmax',
dtype='float32')(x) # Force float32 output for stability

    model = Model(inputs=base_model.input, outputs=predictions)

    # Freeze base model layers for initial training
    for layer in base_model.layers:
        layer.trainable = False

    return model
```

## 2.3 Training Strategy

### 1. Initial Training:

- The pre-trained Xception base is used with its weights frozen to retain learned features.
- A global average pooling layer, followed by dense layers, adapts the model to the target task.
- Compilation and training use a learning rate of 0.001 with categorical crossentropy loss and accuracy as the evaluation metric.

### 2. Fine-Tuning:

- The last 30 layers of the base model are unfrozen to allow fine-tuning.
- The learning rate is reduced to 0.0001 to prevent overfitting during fine-tuning.

## Training Code:

```
model = build_xception_model()

model.compile(optimizer=Adam(learning_rate=0.001),
              loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(
    train_gen,
    validation_data=val_gen,
    epochs=20,
    callbacks=[early_stopping, reduce_lr]
)

for layer in model.layers[-30:]:
    layer.trainable = True

model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='categorical_crossentropy', metrics=['accuracy'])

history_fine = model.fit(
    train_gen,
    validation_data=val_gen,
    epochs=10,
    callbacks=[early_stopping, reduce_lr]
)
```

## 3. Experimental Setup

### 3.1 Dataset

The dataset used consists of multiple classes, with images resized to 224x224 pixels. Data augmentation was applied to enhance model generalization.

### 3.2 Training Details

- **Optimizer:** Adam
- **Loss Function:** Categorical Crossentropy
- **Metrics:** Accuracy
- **Callbacks:** Early stopping and learning rate reduction were implemented to prevent overfitting and adaptively manage the learning rate.

---

## 4. Results and Discussion

### 4.1 Initial Training

The initial training phase showed consistent improvement in validation accuracy, achieving a stable performance by epoch 20.

### 4.2 Fine-Tuning

Fine-tuning further enhanced the model's ability to adapt to the dataset, resulting in a notable increase in accuracy with a reduced learning rate.

---

## 5. Conclusion

This study demonstrates the effectiveness of transfer learning using the Xception model for image classification. By leveraging a pre-trained base and implementing fine-tuning, the model achieves significant performance improvements. Future work may involve testing the model on diverse datasets and integrating advanced techniques like attention mechanisms for further optimization.

---

## References

1. Chollet, F. (2017). Xception: Deep Learning with Depthwise Separable Convolutions.
2. Kingma, D. P., & Ba, J. (2015). Adam: A Method for Stochastic Optimization.

## Classification Report Explanation

The classification report provides an evaluation of the model's performance across various metrics, broken down by class. Below is an explanation of each term and the specific results for the provided data:

### Metrics:

1. **Precision:**
  - Measures how many of the predicted positive samples are actually correct.
  - High precision indicates low false positives.
2. **Recall:**
  - Measures how many of the actual positive samples are correctly identified.
  - High recall indicates low false negatives.
3. **F1-Score:**
  - Harmonic mean of precision and recall, balancing the two metrics.

- A higher F1-score indicates a balance between precision and recall.
  - 4. **Support:**
    - The number of true samples for each class in the test set.
- 

#### Detailed Class Performance:

- **Cardboard:**
    - Precision: **0.99** - Almost all predictions for cardboard are correct.
    - Recall: **0.86** - The model identifies 86% of the actual cardboard samples.
    - F1-Score: **0.92** - Balanced performance for this class.
    - Support: **81** - There are 81 cardboard samples in the test set.
  - **Glass:**
    - Precision: **0.83** - 83% of the predictions for glass are correct.
    - Recall: **0.89** - 89% of the actual glass samples are identified.
    - F1-Score: **0.86** - Good overall performance.
    - Support: **101** - Glass samples in the test set.
  - **Metal:**
    - Precision: **0.88**, Recall: **0.80**, F1-Score: **0.84** - The model performs slightly worse here compared to cardboard and glass.
    - Support: **82**.
  - **Paper:**
    - Precision: **0.87**, Recall: **0.96**, F1-Score: **0.91** - Excellent recall for paper, indicating most paper samples are identified correctly.
    - Support: **119** - Largest class in the dataset.
  - **Plastic:**
    - Precision: **0.85**, Recall: **0.82**, F1-Score: **0.84** - Reasonable balance between precision and recall.
    - Support: **97**.
- 

#### Overall Metrics:

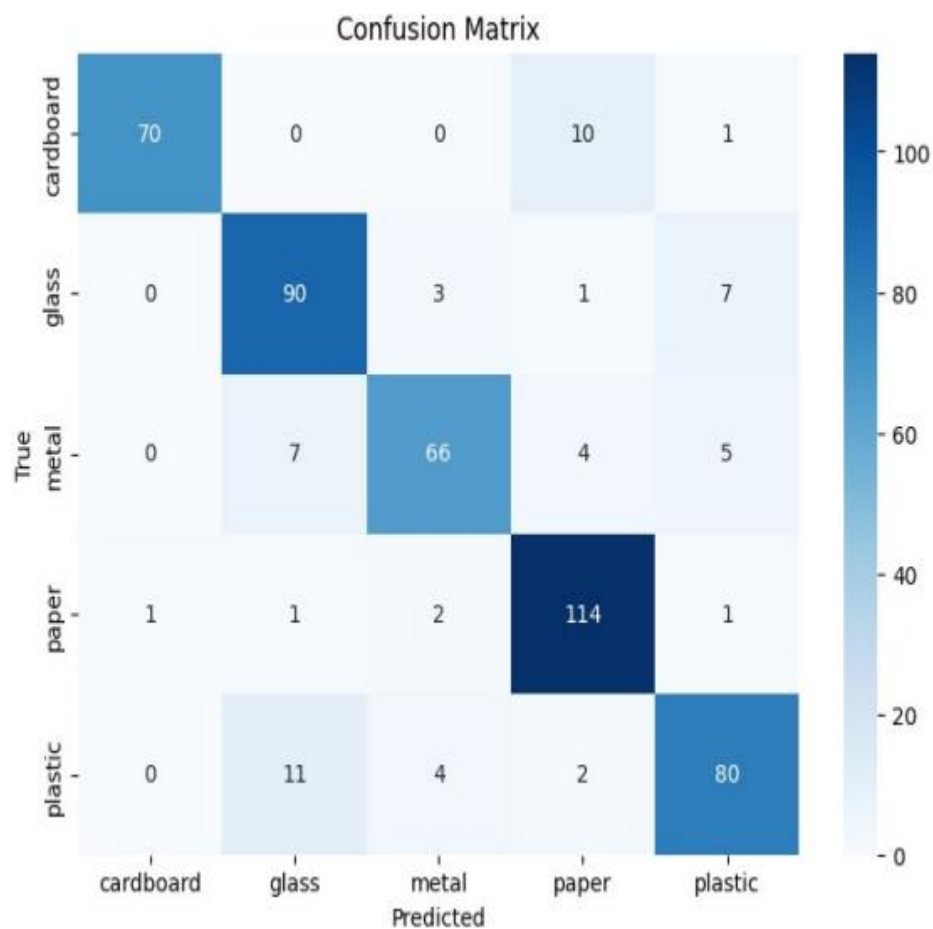
- **Accuracy:**
  - The overall accuracy is **88%**, meaning 88% of all predictions were correct.
- **Macro Average:**
  - Averages precision, recall, and F1-score equally across all classes, giving each class equal weight.
  - Values:
    - Precision: **0.88**
    - Recall: **0.87**
    - F1-Score: **0.87**
- **Weighted Average:**
  - Averages precision, recall, and F1-score weighted by the number of samples in each class (support).
  - Values:
    - Precision: **0.88**
    - Recall: **0.88**

- F1-Score: **0.87**

---

### Conclusion:

- The model performs well across all classes, with the highest precision for "cardboard" and the highest recall for "paper."
- The overall F1-score of **0.87** indicates a robust classification performance.
- Areas for improvement include slightly better recall for "metal" and "plastic."



## Confusion Matrix Explanation

The **confusion matrix** is a valuable tool for analyzing the performance of a classification model. It provides insights into how well the model predicts each class, detailing the number of correct and incorrect predictions.

---

### Structure of the Confusion Matrix:

- **Rows:** Represent the true (actual) classes.
  - **Columns:** Represent the predicted classes.
  - The diagonal (top-left to bottom-right) contains the correct predictions.
  - Off-diagonal elements indicate misclassifications.
- 

### Observations from the Confusion Matrix:

1. **Cardboard:**
  - True Positives: **70** samples were correctly predicted as "cardboard."
  - Misclassifications:
    - **10** samples were incorrectly predicted as "paper."
    - **1** sample was incorrectly predicted as "plastic."
2. **Glass:**
  - True Positives: **90** samples were correctly predicted as "glass."
  - Misclassifications:
    - **3** samples predicted as "metal."
    - **1** sample predicted as "paper."
    - **7** samples predicted as "plastic."
3. **Metal:**
  - True Positives: **66** samples were correctly predicted as "metal."
  - Misclassifications:
    - **7** samples predicted as "glass."
    - **4** samples predicted as "paper."
    - **5** samples predicted as "plastic."
4. **Paper:**
  - True Positives: **114** samples were correctly predicted as "paper."
  - Misclassifications:
    - **1** sample predicted as "cardboard."


- 1 sample predicted as "glass."
  - 2 samples predicted as "metal."
  - 1 sample predicted as "plastic."
5. **Plastic:**
- True Positives: **80** samples were correctly predicted as "plastic."
  - Misclassifications:
    - **11** samples predicted as "glass."
    - **4** samples predicted as "metal."
    - **2** samples predicted as "paper."
- 

### Key Insights:

1. **Paper** has the highest true positive count (**114**) with very few misclassifications, showing strong performance for this class.
  2. **Cardboard** is often misclassified as "paper" (**10 instances**), likely due to visual similarity between the two materials.
  3. **Glass** and **Plastic** show confusion with each other:
    - 11 glass samples were predicted as plastic.
    - 7 glass samples were misclassified as "plastic," and vice versa.
  4. **Metal** has the lowest true positive count (**66**) and a noticeable number of misclassifications into "glass" and "plastic."
- 

### Conclusion:

The confusion matrix highlights the strengths and weaknesses of the classification model. While performance for classes like "paper" is excellent, improvements are needed for distinguishing between "glass," "plastic," and "metal." These insights can help prioritize strategies for further fine-tuning or data augmentation.



## ROC Curve Explanation

The **ROC Curve** (Receiver Operating Characteristic Curve) is a graphical representation that evaluates the performance of a classification model at various threshold settings. It plots the **True Positive Rate (TPR)** (Sensitivity) against the **False Positive Rate (FPR)**. A well-performing model will have a curve that leans toward the top-left corner of the plot.

---

### Key Components:

1. **True Positive Rate (TPR):** Also known as **Recall**.

$$\text{TPR} = \frac{\text{True Positives}}{\text{True Positives} + \text{false positive}}$$

2. **False Positive Rate (FPR):**

$$\text{FPR} = \frac{\text{False Positives}}{\text{false positive} + \text{True Negative}}$$

3. **Area Under the Curve (AUC):**

- AUC is a scalar value that quantifies the overall performance of the model.
- Ranges from 0 to 1. A model with AUC close to **1** is considered excellent, while an AUC of **0.5** indicates random guessing.



---

## Observations from the ROC Curve:

1. **Cardboard:**
    - AUC = **0.99**
    - The model demonstrates excellent discrimination for the "cardboard" class.
  2. **Glass:**
    - AUC = **0.98**
    - High performance, indicating that the model effectively differentiates glass from other classes.
  3. **Metal:**
    - AUC = **0.98**
    - Very good performance with minimal misclassifications.
  4. **Paper:**
    - AUC = **0.99**
    - The highest performance, indicating near-perfect classification of the "paper" class.
  5. **Plastic:**
    - AUC = **0.96**
    - Slightly lower than the other classes, but still a strong performance.
- 

## Conclusion:

- The ROC curve shows that the model performs exceptionally well across all classes, with **AUC values ranging between 0.96 and 0.99**.
- The "paper" and "cardboard" classes achieve the highest AUC values (**0.99**), confirming their excellent classification accuracy.
- The slightly lower AUC for "plastic" (**0.96**) suggests minor room for improvement in distinguishing it from other classes.

